

CSC 411

Computer Organization (Spring 2024) Lecture 10: RISC-V

Prof. Marco Alvarez, University of Rhode Island

The RISC-V instruction set

- We will learn about ISA design by learning RISC-V
 - modern and full featured RISC ISA
 - open source, free, simple, extensible, support heterogeneous and parallel systems, supports 32-bit and 64-bit variants
 - supports (but does not require) IEEE 754
- Developed at UC Berkeley as open ISA (~2010)
 - by 2020, more than 200 companies are members of RISC-V International (riscv.org)
- Similar ISAs have a large share of embedded core market
 - applications in consumer electronics, network/storage equipment, cameras, printers, ...



RISC-V base and extensions

Name	Description	Version	Status ^[A]	Instruction count
Base				
RVWM0	Weak Memory Ordering	2.0	Ratified	
RV32I	Base Integer Instruction Set, 32-bit	2.1	Ratified	40
RV32E	Base Integer Instruction Set (embedded), 32-bit, 16 registers	2.0	Ratified	40
RV64I	Base Integer Instruction Set, 64-bit	2.1	Ratified	15
RV64E	Base Integer Instruction Set (embedded), 64-bit	2.0	Ratified	
RV128I	Base Integer Instruction Set, 128-bit	1.7	Open	15
Extension				
M	Standard Extension for Integer Multiplication and Division	2.0	Ratified	8 (RV32) 13 (RV64)
A	Standard Extension for Atomic Instructions	2.1	Ratified	11 (RV32) 22 (RV64)
F	Standard Extension for Single-Precision Floating-Point	2.2	Ratified	26 (RV32) 30 (RV64)
D	Standard Extension for Double-Precision Floating-Point	2.2	Ratified	26 (RV32) 32 (RV64)

and many others ... <https://en.wikipedia.org/wiki/RISC-V>


Instruction set

- Assembly language and ISAs
 - assembly language acts as a bridge between the human-readable world and the binary world of machine code
 - each assembly language is specific to a particular ISA and reflects its instruction set and architecture
 - each line of assembly code represents one instruction for the computer
- Assembly operands are **registers**
 - RISC-V operations can only be performed on registers

```
add x1, x2, x3 // x1 <- x2 + x3
```

Registers

- Built directly into the hardware
 - limited number of registers available (**register file**)
 - assembly code must be very carefully put together to efficiently use them (mostly done by **compilers**)
 - very fast, located within the CPU (0-1 cycles)
 - on a 3 GHz CPU would have an access time of approximately 0.33 ns
- Registers have no type
 - contents are just bits, the **“operation”** determines how the bits are interpreted



registers are extremely fast with a very limited capacity

credit: Berkeley CS61C Great Ideas in Computer Architecture

RISC-V registers

- RISC-V has 32 integer registers (or 16 in the embedded variant)
 - a larger number can lead to an increase in clock cycle time, as the the physical distance between them and other components like the ALU increases
 - numbered from **x0 to x31** and can be referenced by number or name
 - additional 32 floating-point registers when the floating-point extension is implemented
 - each register is 32-bit wide (RV32 variant)
- 32-bit sequences are called a **word** in RV32 and 64-bits sequences a **doubleword**
- RISC-V is a load-store architecture, instructions address only registers,
 - load and store instructions convey data to and from memory
- No instructions exist to save and restore multiple registers

RISC-V registers

Register name	Symbolic name	Description	Saved by
32 integer registers			
x0	Zero	Always zero	
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	
x4	tp	Thread pointer	
x5	t0	Temporary / alternate return address	Caller
x6-7	t1-2	Temporary	Caller
x8	s0/lp	Saved register / frame pointer	Callee
x9	s1	Saved register	Callee
x10-11	a0-1	Function argument / return value	Caller
x12-17	a2-7	Function argument	Caller
x18-27	s2-11	Saved register	Callee
x28-31	t3-6	Temporary	Caller

<https://en.wikipedia.org/wiki/RISC-V>

RISC-V registers

32 floating-point extension registers			
f0-7	ft0-7	Floating-point temporaries	Caller
f8-9	fs0-1	Floating-point saved registers	Callee
f10-11	fa0-1	Floating-point arguments/return values	Caller
f12-17	fa2-7	Floating-point arguments	Caller
f18-27	fs2-11	Floating-point saved registers	Callee
f28-31	ft8-11	Floating-point temporaries	Caller

<https://en.wikipedia.org/wiki/RISC-V>

Arithmetic operations

- Require three register **operands**
 - all arithmetic operations have this rigid form: opname, destination, source 1, source 2

opname rd, rs1, rs2

	Addition	Subtraction
c	a = b + c;	a = b - c;
RISC-V	add x1, x2, x3	sub x1, x2, x3

Arithmetic operations

- How would you translate the following C code?
 - single line of C may convert into multiple lines in assembly
 - some sequences are shorter than others or may use less “temporary” registers

```
// assume these variables are mapped to  
// x5, x1, x2, x3, x4 respectively  
a = b + c + d + e;
```

Practice

- Translate the following C code into assembly
 - if needed, use temporary registers x6, x7

```
// assume these variables are mapped to  
// x5, x1, x2, x3, x4 respectively  
f = (g + h) - (i + j);
```

- optimize the C code to minimize register usage
 - good compilers do it all the time

In fact, if the variables are floating-point values, different sequences of instructions may produce slightly different results. Floating-point operations are not necessarily associative or commutative ...

Immediates

- Immediates are just numerical constants
 - an immediate operand avoids a **load** instruction
- Immediate instructions
 - instructions where constant values may be specified
- No subtract immediate instruction
 - there are **add** and **sub** instructions but only **addi** for immediate operands (just use a negative constant)
 - e.g.
`addi x5, x6, 5`
`addi x5, x6, -10`

Register **zero**

- Register **x0** is hardwired to the 0
 - **cannot be overwritten**: store to the zero register has no effect, and a read always provides 0
- Useful for common operations
 - e.g., move between registers, assigning constants to registers